

ImageBIND

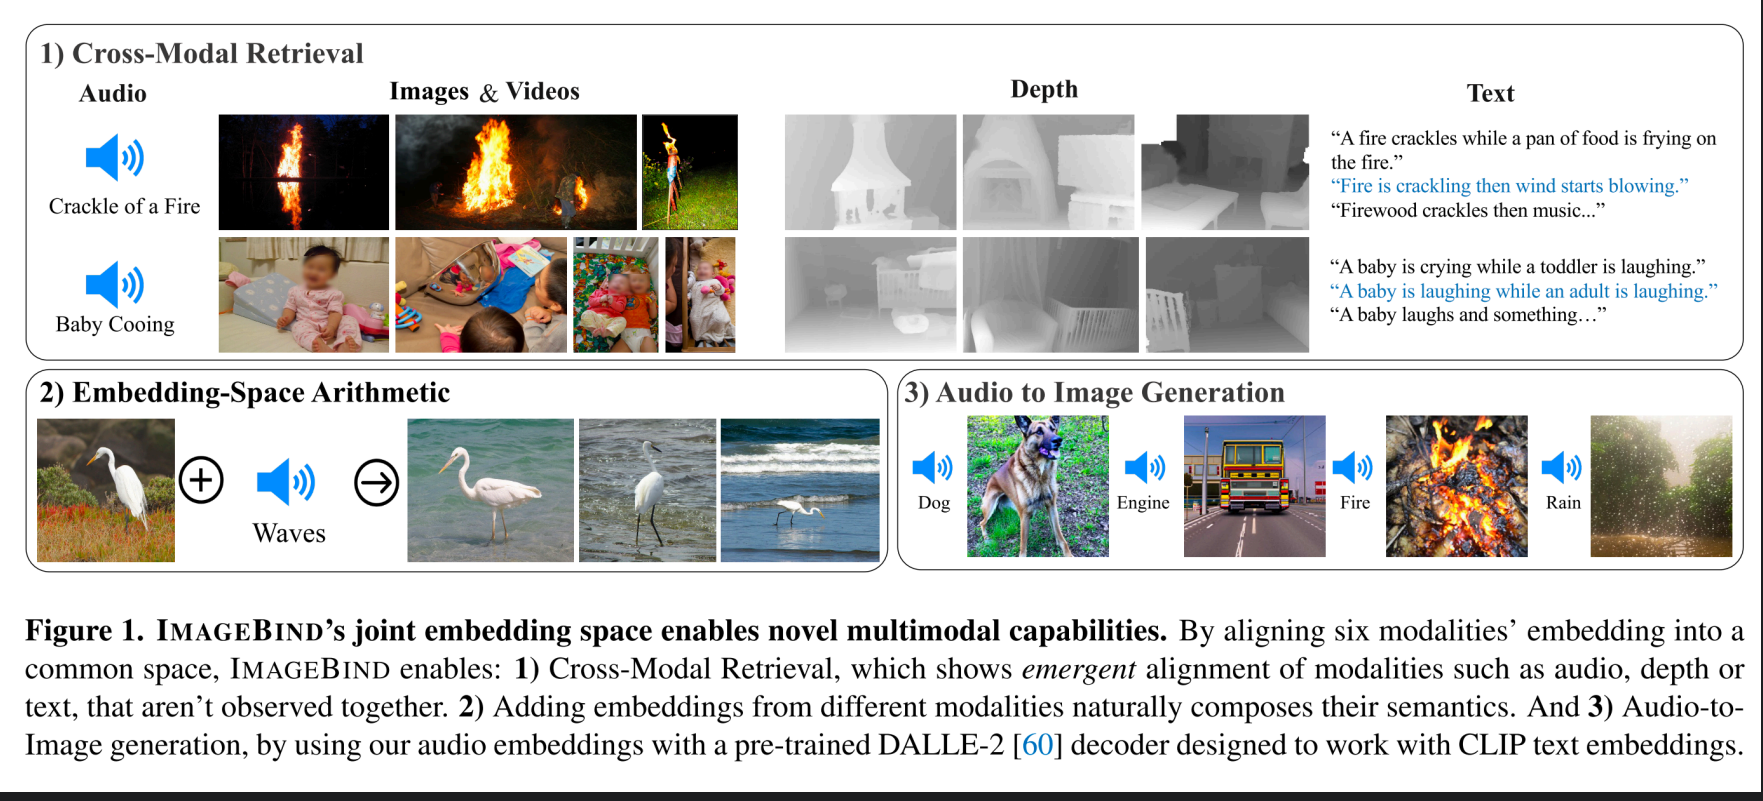


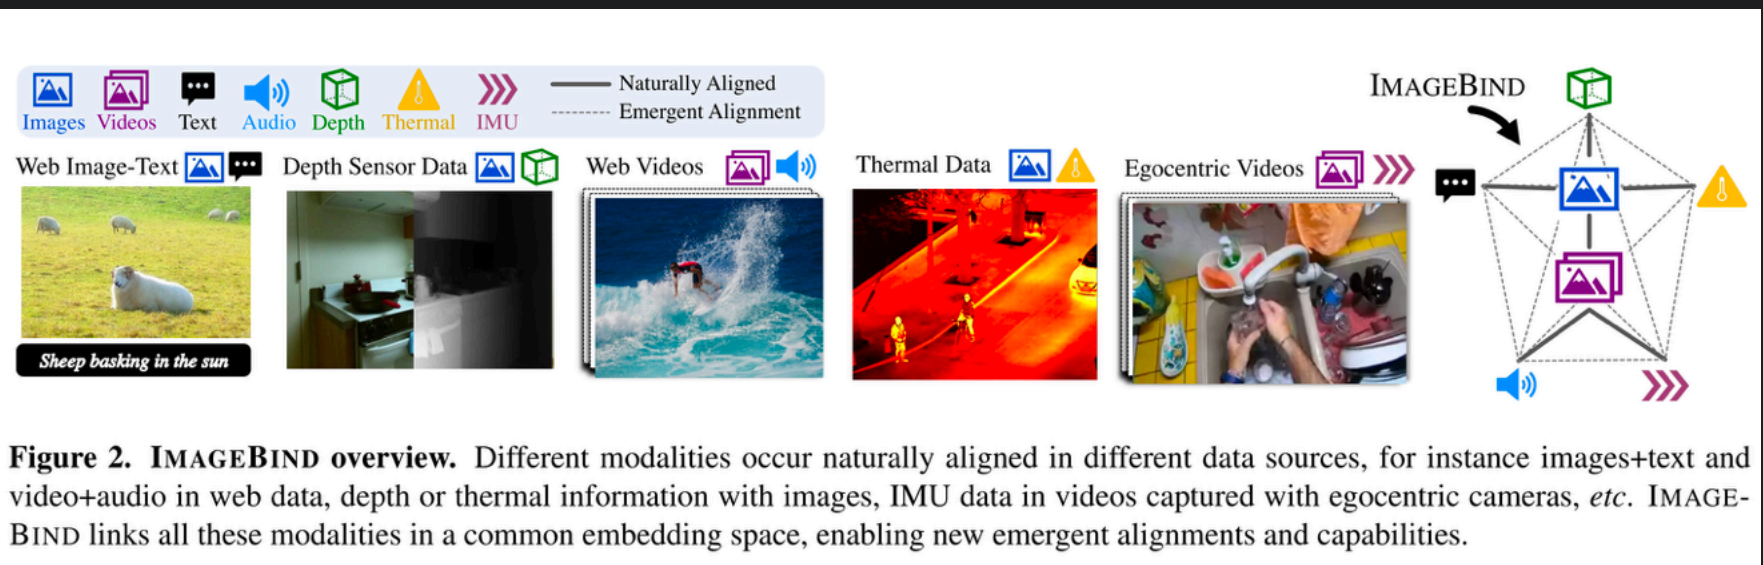
Figure 1. IMAGEBIND’s joint embedding space enables novel multimodal capabilities. By aligning six modalities’ embedding into a common space, IMAGEBIND enables: **1)** Cross-Modal Retrieval, which shows *emergent* alignment of modalities such as audio, depth or text, that aren’t observed together. **2)** Adding embeddings from different modalities naturally composes their semantics. And **3)** Audio-to-Image generation, by using our audio embeddings with a pre-trained DALLE-2 [60] decoder designed to work with CLIP text embeddings.

Abstract

1. An approach to learn a joint embedding across six different modalities – images, text, audio, depth, thermal, and IMU data.
2. All combinations of paired data are not necessary to train such a joint embedding, and **only image-paired data** is sufficient to bind the modalities together.
3. IMAGEBIND can leverage recent large scale vision–language models, and extends their **zero-shot capabilities** to new modalities just by using their natural pairing with images.

Method

1. we present IMAGEBIND, which learns a single shared representation space by leveraging multiple types of image–paired data. Just aligning each modality’s embedding to image embeddings leads to an emergent alignment across all of the modalities.
2. We use large–scale image–text paired data (CLIP) along with naturally paired ‘self–supervised’ data across four new modalities – audio, depth, thermal, and Inertial Measurement Unit (IMU) readings – and show strong emergent zero–shot classification and retrieval performance on tasks for each of these modalities.



InfoNCE

$$L_{\mathcal{I}, \mathcal{M}} = -\log \frac{\exp(\mathbf{q}_i^\top \mathbf{k}_i / \tau)}{\exp(\mathbf{q}_i^\top \mathbf{k}_i / \tau) + \sum_{j \neq i} \exp(\mathbf{q}_i^\top \mathbf{k}_j / \tau)}, \quad (1)$$

NCE: Noise–Contrastive Estimation

we use a symmetric loss $L_{\mathcal{I}, \mathcal{M}} + L_{\mathcal{M}, \mathcal{I}}$.

Implamentation details

1. Use Transformer for all the modality encoder
2. Vision Transformer (ViT) for images, videos (2 frame from 2 sec video, inflate ViT patch projection), depth(1D image) & therimal images(1D images)
3. Audio: 2sec, 16khz into spectrograms using 128 mel–spectrogram bins. As 2D signal, patch size of 16, stride 10, ViT
4. IMU: IMU signal consisting of accelerometer and gyroscope measurements across the X, Y, and Z axes. We use 5 second clips resulting in 2K time step IMU readings which are projected using a 1D convolution with a kernel size of 8. Then use transformer.
5. Text: CLIP

Text and image encoder the same as CLIP!